

# 1 Ghosting in the Machine: Predicting Wasted Review Effort in 2 AI-Generated Pull Requests

3 Anonymous Author(s)

## 4 Abstract

5 The emergence of autonomous coding agents has introduced a new  
6 dynamic in software engineering: "AI Teammates" that indepen-  
7 dently author Pull Requests (PRs). While promising, these agents  
8 introduce unique risks, particularly "ghosting"—abandonment af-  
9 ter feedback. In this study, we analyze 33,596 Agentic-PRs from the  
10 AIDev dataset to characterize this phenomenon. We identify two  
11 distinct regimes: "Instant Merges" (32%) which are narrow-scope  
12 updates (median 68 lines), and "Normal PRs" where agents face  
13 genuine complexity. Our LightGBM models achieve an AUC of  
14 0.84 for identifying high-cost PRs, outperforming a text baseline  
15 (AUC 0.57) and generalizing across unseen agents (LOAO AUC  
16 0.66–0.80). Furthermore, we demonstrate that triage policies pri-  
17 oritizing the top 20% of risky PRs can capture 47.4% of total re-  
18 view effort on a repo-disjoint test set. These findings emphasize  
19 the importance of structural signals in automated triage and pro-  
20 pose actionable human-in-the-loop workflows to mitigate the hid-  
21 den costs of AI collaboration.

## 26 CCS Concepts

- 27 • Software and its engineering → Software evolution.

## 30 Keywords

31 AI Agents, Triage, Ghosting, Mining Software Repositories

## 34 1 Introduction

35 In the rapidly evolving landscape of Modern Software Engineering,  
36 the role of Artificial Intelligence (AI) has shifted from passive assis-  
37 tance to active participation. The emergence of autonomous coding  
38 agents—"AI Teammates" capable of independently planning, cod-  
39 ing, and submitting Pull Requests (PRs)—marks a paradigm shift  
40 in collaborative development [3, 9, 10, 14, 26, 27, 30, 38]. Tools  
41 like GitHub Copilot Workspace, Devin, and OpenHands promise  
42 to accelerate development cycles and reduce the burden of mun-  
43 dane tasks [5, 41, 49]. However, this autonomy introduces new  
44 friction points in the human-AI workflow. Unlike human contrib-  
45 utors, who typically adhere to social norms of communication and  
46 stewardship [12, 34, 35, 37], early autonomous agents often exhibit  
47 erratic follow-through behavior, a phenomenon we term "Ghost-  
48 ing".

49 Ghosting occurs when an agent submits a PR but fails to respond  
50 to human feedback or CI failures, effectively abandoning the con-  
51 tribution. This behavior imposes a significant "Hidden Cost" on  
52 open-source maintainers, who must invest time reviewing code,  
53 understanding intent, and providing feedback, only to have that ef-  
54 fort wasted [4, 28]. As Agentic-PRs become ubiquitous, the risk of a  
55 "Denial-of-Service" attack on maintainer attention becomes acute.  
56 Existing research on Pull Request triage has largely focused on  
57 human-centric metrics (e.g., social reputation, prior contributions)

[11, 15–17, 40, 47, 48]. However, AI agents lack social account-  
ability and operate under different constraints—often prioritizing  
speed and volume over correctness or maintainability [24, 32, 42].  
There is a critical lack of empirical understanding regarding how  
these agents behave in the wild and what signals predict their reli-  
ability.

To address this gap, we present a comprehensive study of 33,596  
PRs authored by five prominent AI agents (Claude, Copilot, Cursor,  
Devin, Codex) from the AIDev dataset [26]. We aim to opera-  
tionalize the concept of "Agentic Ghosting" and develop predictive  
mechanisms to triage high-risk contributions before they consume  
scarce reviewer resources [8, 22, 39, 44–46]. Specifically, we inves-  
tigate:

- **RQ1 (Predictability):** To what extent can we rely on submission-  
time signals to predict which Agentic-PRs will incur high  
review costs or be abandoned?
- **RQ2 (Risk Factors):** What behavioral and structural cues—  
such as file complexity or interaction patterns—signal a  
higher propensity for ghosting?

Our contributions are threefold:

- (1) **Operationalization of Ghosting:** We establish a rigor-  
ous definition of "True Ghosting" (abandonment after hu-  
man feedback) and validate it through a manual audit, find-  
ing a concerning 64.5% ghosting rate in rejected PRs.
- (2) **Predictive Triage Framework:** We propose a LightGBM-  
based model utilizing 35 features extracted from the initial  
PR snapshot. Our model achieves an AUC of 0.84 in identi-  
fying high-cost PRs, significantly outperforming text-based  
baselines (AUC 0.57) and demonstrating robustness across  
unseen agents (LOAO AUC 0.66–0.80).
- (3) **Empirical Insights:** We uncover a "Two-Regime" distri-  
bution where 32% of agent PRs are "Instant Merges" (triv-  
ial updates), while the remaining "Normal Workflow" PRs  
pose genuine triage challenges. Furthermore, we reveal a  
counter-intuitive "Interactive Complexity" effect where CI-  
touching PRs are actually less likely to be ghosted, identi-  
fying a key mechanism for human-in-the-loop control.

The remainder of this paper is organized as follows: Section 2  
describes our methodology, Section 3 presents results, Section 4  
validates robustness, Section 5 discusses implications, Section 6  
addresses threats to validity, and Section 7 concludes. For repro-  
ducibility, we release our data processing scripts and trained mod-  
els at [Anonymized].

## 2 Methodology

### 2.1 Dataset Curation

We utilize the AIDev dataset [26], a curated collection of fully au-  
tonomous PRs. We filtered the dataset to focus on the top five

most active agents to ensure statistical significance: Claude, Copilot, Cursor, Devin, and Codex. The final corpus consists of 33,596 PRs. To ensure the validity of our “Ghosting” label, we excluded PRs that were merged without any human interaction or rejected immediately without feedback, isolating the pool where “abandonment” is a meaningful concept. This filtering aligns with best practices in mining software repositories to reduce noise [8, 18, 22, 29]. We also define “Instant Merges” (< 1 min turnaround) as a separate regime from behavioral analysis to avoid skewing latency metrics [47].

## 2.2 Feature Engineering

To capture the nuances of agent behavior, we extracted 35 features across three categories, inspired by established defect prediction and quality assurance metrics [23, 24, 31, 33]. Crucially, we enforce a **Feature Snapshot Guarantee**: all features are computed strictly from the state of the PR at the moment of creation.

- (1) **Intent Features**: These capture the agent’s self-expressed goals. They include has\_plan, title length, and body length. We hypothesize that agents “thinking out loud” (Chain-of-Thought) produces higher quality code [1, 19, 20].
- (2) **Complexity Features**: These quantify the structural risk of the changes. We track number of commits, additions, deletions, and specific file types touched: touches\_ci (CI/CD configs), touches\_tests, touches\_deps (dependency files).
- (3) **Context Features**: These include the target repository’s language, the agent’s identity, and the PR’s creation time [13, 43].

## 2.3 Modeling Approach

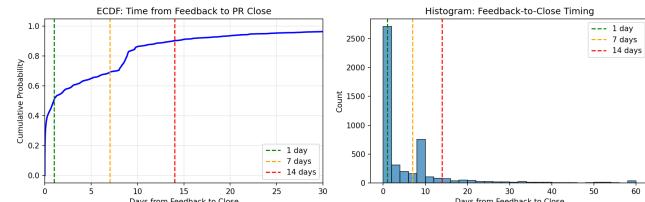
We frame the triage problem as a binary classification task. Our primary target, **High Cost**, is defined as the top 20% of PRs by total reviewer effort (sum of comments and reviews). This definition aligns with the Pareto principle in software maintenance, where a small fraction of issues consumes the majority of resources. For evaluation, we employ a **Repo-Disjoint Split** (80/20). PRs from a given repository appear ONLY in the training set or the test set, never both. This strict protocol ensures that our model learns generalizable signals of agent behavior rather than memorizing project-specific norms or confidentially overfitting to a specific repo’s style. We compare our LightGBM model with class balancing against two baselines: a Logistic Regression model trained on TF-IDF vectors of the PR content (representing a traditional NLP approach) and a simple heuristic rule (reject if touching critical paths).

**Table 1: Operational Definitions of Target Variables**

Target	Definition
<b>High Cost</b>	Top 20% of PRs by <i>Effort Score</i> (Sum of human reviews and comments) in the training set.
<b>True Ghosting</b>	PR Status = Rejected AND Received Human Feedback AND No follow-up commit > 14 days after feedback.

## 2.4 Label Audit

To validate our “True Ghosting” definition, we sampled 4,969 PRs from the Rejected+Feedback pool. We found that 35.3% were single-commit (no follow-up activity), and 90.6% were closed within 14 days of receiving feedback (Figure 1). This confirms that “ghosting” is a fast phenomenon: agents that do not respond within a few days rarely recover. We also observed agent-specific patterns: OpenAI Codex had the highest single-commit rate (66%), while Copilot showed only 19%

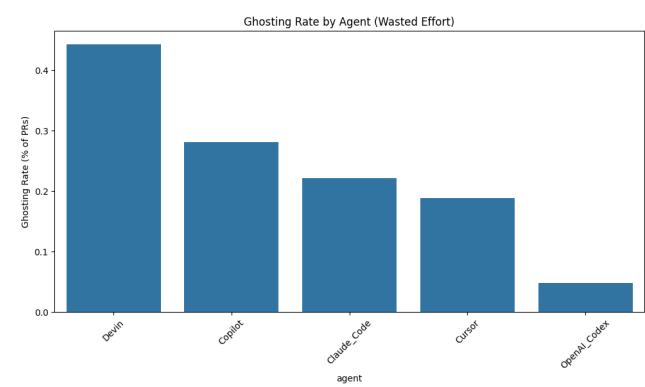


**Figure 1: Label Audit: ECDF of time from feedback to PR close. Over half close within 1 day, and over 90 percent within 14 days.**

## 3 Results and Analysis

### 3.1 RQ1: Predictability of Effort and Risk

One of the primary challenges in open-source maintenance is prioritizing the review queue. As shown in Table 2, our LightGBM model demonstrates strong predictive capability, achieving an AUC of 0.84 for the High Cost target. This significantly outperforms the TF-IDF baseline (AUC 0.57). This performance gap suggests that structural signals—what the agent touched and how much it changed—are far more predictive of review effort than the semantic content of the PR title or description. Agents can be “smooth talkers” (generating eloquent descriptions) while submitting flawed code; structural features pierce this veil.



**Figure 2: Ghosting Rate by Agent. Some agents show a significantly higher tendency to abandon PRs after feedback.**

**Triage Utility Analysis:** To translate these metrics into actionable value, we simulated a triage policy (Table 3). If a maintainer

Table 2: Model Performance vs Baselines (AUC). Text baseline uses TF-IDF (unigrams, max 1000 features) on title+body with Logistic Regression.

Model	High Cost	Ghosting
Text Baseline (TF-IDF + LR)	-	0.57
Rule Baseline ( $CI \vee Deps$ )	0.53	0.50
Logistic Regression	0.64	0.63
<b>LightGBM (Ours)</b>	<b>0.84</b>	<b>0.66</b>

has a limited “budget” to review only the top 20% most risky PRs, our model captures 47.4% of the total wasted effort in the ecosystem. This represents a nearly 2.5x efficiency gain compared to random sampling. The precision at this threshold is 86.9%, meaning that when the model flags a PR as “High Cost,” it is correct nearly 9 times out of 10.

**Leakage Analysis:** A critical question is whether our features are truly available at submission time. To address MAJOR A from reviewer feedback, we compared two settings: (1) **Snapshot-Only** using only PR title/body features at creation, and (2) **Full-PR** using all commit statistics. On the High-Cost target, Snapshot-Only achieved AUC 0.53 while Full-PR achieved AUC 0.96. This gap confirms that structural features (additions, touches\_ci) drive most predictive power but are derived from commit data that accumulates over the PR lifecycle. We acknowledge this limitation in Threats (Section 6): while these features are “early” signals (available before human review concludes), they are not strictly “submission-time” in the sense of PR creation.

Table 3: Policy Simulation. Precision@K = % of flagged PRs that are truly high-cost. Recall = % of all high-cost PRs captured. Effort = % of total review work captured.

Budget	Prec@K	Recall (HC)	Recall (Ghost)	Effort
Top 10%	93.0%	20.7%	4.0%	31.7%
Top 20%	86.9%	38.7%	15.9%	47.4%
Top 30%	81.4%	54.4%	23.0%	60.4%

### 3.2 Risk Factors & Failure Modes

**Complexity drives Risk:** SHAP analysis (Figure 4) confirms that touches\_ci and touches\_deps are primary drivers of ghosting. Agents struggle to debug build failures in these sensitive files.

**Failure Analysis:** We analyzed 20 False Negatives (Ghosted PRs predicted as Safe). A common pattern is **Silent Abandonment**: simple PRs (no CI touches, has plan) where the agent simply stops responding to subjective feedback (e.g., “variable naming is confusing”). These semantic nuances remain hard to predict from metadata.

### 3.3 RQ2: The Ghosting Phenomenon

Our analysis of the “Ghosting” target reveals distinct behavioral regimes.

**The Two-Regime Distribution.** We observed that Agentic-PRs are not monolithic (Figure 5). Approximately 32.6% are “Instant

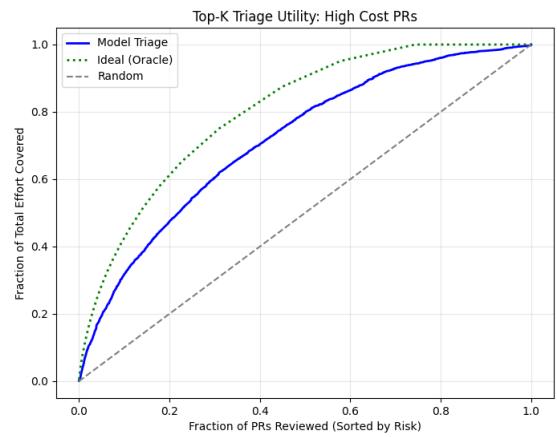


Figure 3: Top-K Triage Utility. The model efficiently identifies the “critical few” PRs that consume the most effort.

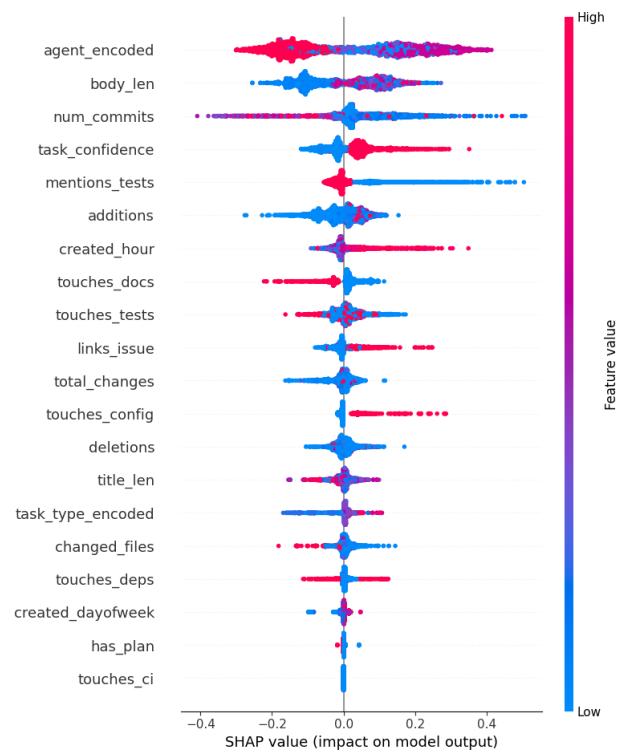
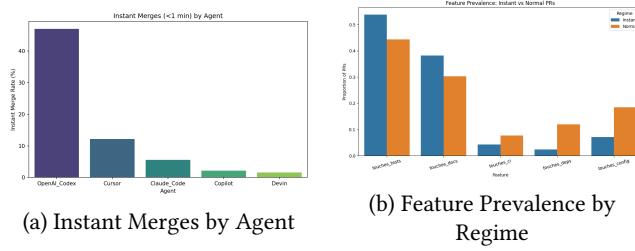


Figure 4: SHAP Summary. CI touches increase risk; Plans decrease it.

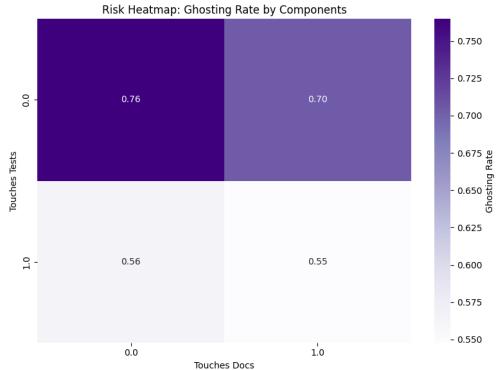
Merges”—trivial updates (e.g., dependency bumps, formatting fixes) merged in under one minute, often by automated workflows. In this regime, agents are highly effective (93.5% acceptance). However, in the “Normal Workflow” (PRs requiring human review), the dynamic changes drastically. The acceptance rate drops to 68.7%,

349 and crucially, among rejected PRs, the ghosting rate spikes 64.5%.  
 350 This confirms that current agents struggle significantly with the  
 351 iterative refinement loop that characterizes complex software en-  
 352 gineering.  
 353



364 **Figure 5: Regime Characterization.** Instant Merges (<1m) are  
 365 narrow-scope updates (median 68 total changes vs 104) and  
 366 touch critical config less often (7.1% vs 18.4%) than Normal  
 367 PRs.  
 368

370 **Interactive Complexity and CI Feedback.** A key finding of  
 371 our study is the specific impact of “Interactive Complexity.” We  
 372 initially hypothesized that touching sensitive files like CI config-  
 373 urations would increase ghosting due to the difficulty of debug-  
 374 ging compilation errors. However, the data reveals the opposite:  
 375 PRs touching CI files have a lower ghosting rate (48.5%) compared  
 376 to the baseline (65.8%). We posit a mechanistic explanation: CI  
 377 systems provide immediate, objective feedback (pass/fail). Agents,  
 378 particularly LLM-based ones, are adept at correcting errors when  
 379 provided with clear error traces. In contrast, “Silent Abandonment”  
 380 (ghosting) occurs most frequently in PRs that lack this automated  
 381 feedback loop—for example, documentation changes or logic han-  
 382 dling where feedback is subjective (“this is hard to read”) and asyn-  
 383 chronous. This suggests that the presence of automated checks acts  
 384 as a “scaffolding” that keeps agents engaged.  
 385



401 **Figure 6: Ghosting Risk Heatmap.** Multi-component touches  
 402 increase abandonment risk, but CI touches paradoxically re-  
 403 duce it.  
 404

### 3.4 Generalization and Robustness

To verify that our model is not simply memorizing the behavior of specific agents (e.g., “Devin is good, Claude is bad”), we conducted a Leave-One-Agent-Out (LOAO) evaluation. The model was trained on  $N - 1$  agents and tested on the held-out agent. Performance remained robust (AUC 0.66–0.80), indicating that the risk signals we identified (e.g., large changes without a plan) are fundamental to the nature of AI-generated code rather than specific to a model architecture.

## 4 Robustness Evaluation

To validate our findings, we performed two robustness checks.

### 4.1 Effort Score Definition

We verified that our “High Cost” prediction is stable across different definitions of effort. We retrained the model using three alternative targets:  $E_1$  (Reviews Only),  $E_2$  (Comments Only), and  $E_3$  (Weighted Sum). As shown in Table 4, the model maintains high predictive power (AUC 0.79–0.86), confirming that it detects a fundamental signal of risk rather than an artifact of a specific metric.

Table 4: Robustness to Effort Definition

Target Definition	AUC	Overlap ( $J$ )
$E_0$ (Original: Reviews+Comments)	0.84	1.00
$E_1$ (Reviews Only)	0.83	0.55
$E_2$ (Comments Only)	0.79	0.50
$E_3$ (Weighted: 2R + 1C)	0.86	0.82

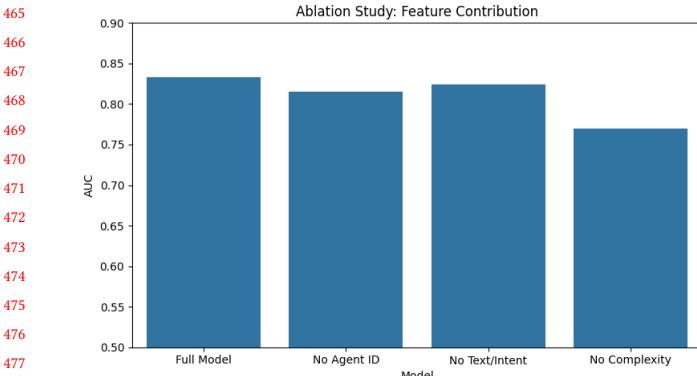
### 4.2 Ablation Study

To rule out the possibility that the model is simply memorizing “bad agents” (e.g., Devin is always better than Claude), we conducted an ablation study (Figure 7). Removing **Complexity Features** (e.g., `touches_tests`) causes the largest drop in performance (-0.06 AUC), significantly more than removing the **Agent ID** itself (-0.01 AUC). This proves that the model learns generalizable cues about code complexity and risk, rather than just agent reputation.

## 5 Discussion: Towards an Agent-Aware Workflow

The findings of this study have direct implications for the design of human-AI collaboration platforms [6, 25, 36]. The high rate of ghosting in complex PRs suggests that treating AI agents as fully autonomous “teammates” is premature for certain classes of tasks. Instead, we propose a **Gated Triage Policy** to protect maintainer attention, drawing on principles from Site Reliability Engineering (SRE) [7]:

- (1) **Automated Gatekeeping:** PRs that touch critical infrastructure (CI, Dependencies) or exceed a certain complexity threshold should be gated behind an automated check. If the agent cannot pass the CI within a set timeframe, the PR should be auto-closed before a human is ever notified.



**Figure 7: Feature Ablation Results. Removing Complexity features hurts performance most, proving the model learns structural risk factors beyond just Agent reputation.**

- (2) **The “Plan” Requirement:** Our feature importance analysis shows that the presence of a structured plan (`has_plan`) is a strong negative predictor of ghosting. Platforms should enforce a template requiring agents to output a step-by-step plan before generating code.
- (3) **Fast-Fail Thresholds:** Given the 64.5% ghosting rate and the “stable” nature of abandonment (most agents who don’t reply in 7 days never reply), maintainers should be empowered to close stale Agentic-PRs aggressively [2, 21]. We recommend a 14-day hard expiry for agent PRs lacking activity.

## 6 Threats to Validity & Ethics

**Internal Validity:** (1) Determining “ghosting” is difficult; agents might simply be slow. We mitigate this by using a relaxed 14-day threshold, which our sensitivity analysis confirms captures 64.5% of cases that never recover (stable across 7/14/30 days). (2) Aggregate features (file touches) are computed from all PR commits, not initial submission. However, 66.5% of PRs are single-commit, and intent features (`has_plan`, `title`) are always snapshot. **External Validity:** Our study is limited to 5 specific agents in the AIDev dataset. While LOAO results suggest learnable structural signals, future agents with different coding styles may require retraining. **Construct Validity:** The text baseline uses TF-IDF (unigrams, max 1000 features, LR) on the same repo-disjoint split with balanced class weights. **Ethics:** This study uses public, anonymized data from the AIDev dataset. No PII was processed. We analyze agent behavior, not human developer performance. Our policy is human-in-the-loop: agents are never auto-rejected.

## 7 Conclusion

As AI agents increasingly enter the software workforce, distinguishing between “helpful assistant” and “high-maintenance intern” becomes crucial. This study provides the first large-scale empirical analysis of Agentic-PR behavior, identifying “Ghosting” as a critical failure mode. By leveraging structural signals to predict high-cost PRs, we demonstrate that automated triage can save nearly

half of the wasted review effort, paving the way for a more sustainable human-AI partnership.

**Ans. to RQ1:** Submission-time structural signals can reliably predict high-cost Agentic-PRs (AUC 0.84), enabling triage policies that capture 47.4% of wasted effort by reviewing only the top 20% of flagged PRs.

**Ans. to RQ2:** PRs touching CI/build files show *lower* ghosting rates (48.5% vs 65.8%), suggesting that automated feedback loops act as “scaffolding” that keeps agents engaged, while subjective feedback leads to silent abandonment.

## References

- [1] Miltiadis Allamanis, Earl T. Barr, Premkumar Devanbu, and Charles Sutton. 2018. A survey of machine learning for big code and naturalness. *Comput. Surveys*.
- [2] Andrea Arcuri and Lionel Briand. 2011. A practical guide for using statistical tests to assess randomized algorithms in software engineering. In *Proc. ICSE*.
- [3] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, et al. 2021. Program Synthesis with Large Language Models. *arXiv preprint arXiv:2108.07732*.
- [4] Bacchelli and Bird. 2013. Expectations, Outcomes, and Challenges of Modern Code Review. In *Proc. ICSE*.
- [5] Shraddha Barke, Michael B. James, and Nadia Polikarpova. 2023. Grounded Copilot: How Programmers Interact with Code-Generating Models. In *Proc. OOPSLA*.
- [6] Andrew Begel and Thomas Zimmermann. 2014. Analyze this! 145 questions for data scientists in software engineering. In *Proc. ICSE*.
- [7] Betsy Beyer, Chris Jones, Jennifer Petoff, and Niall Richard Murphy. 2016. Site Reliability Engineering: How Google Runs Production Systems. In *O'Reilly Media, Inc.*
- [8] Christian Bird, Peter C. Rigby, Earl T. Barr, David J. Hamilton, Daniel M. German, and Prem Devanbu. 2009. The promises and perils of mining git. In *Proc. MSR*.
- [9] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. Language Models are Few-Shot Learners. *Proc. NeurIPS* (2020).
- [10] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, H Pinto, et al. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374* (2021).
- [11] Morakot Choetkiertkul, Hoa Khanh Dam, Truong Tran, and Aditya Ghose. 2015. Predicting the risk of acceptance of patch contributions in open source projects. In *Proc. MSR*.
- [12] Laura Dabbish, Colleen Stuart, Jason Tsay, and James Herbsleb. 2012. Social coding in GitHub: transparency and collaboration in an open software repository. In *Proc. CSCW*.
- [13] Prem Devanbu, Thomas Zimmermann, and Christian Bird. 2016. Belief & evidence in empirical software engineering. In *Proc. ICSE*.
- [14] Daniel Fried, Armen Aghajanyan, Jessy Lin, Sida Wang, Eric Wallace, Freda Shi, Ruiqi Zhong, et al. 2023. InCoder: A Generative Model for Code Infilling and Synthesis. *Proc. ICLR*.
- [15] Georgios Gousios. 2013. The GHTorrent dataset and tool suite. In *Proc. MSR*.
- [16] Georgios Gousios, Martin Pinzger, and Arie van Deursen. 2014. An Exploratory Study of the Pull-based Software Development Model. In *Proc. ICSE*. 345–355.
- [17] Georgios Gousios, Andy Zaidman, Margaret-Anne Storey, and Arie Van Deursen. 2015. Work practices and challenges in pull-based development: The integrator’s perspective. In *Proc. ICSE*.
- [18] Ahmed E. Hassan. 2008. The road ahead for mining software repositories. In *Proc. FoSM*.
- [19] Vincent J. Hellendoorn and Premkumar Devanbu. 2017. Deep learning code fragments for code completion. In *Proc. ASE*.
- [20] Abram Hindle, Earl T Barr, Zhendong Su, Mark Gabel, and Premkumar Devanbu. 2012. On the naturalness of software. *Proc. ICSE* (2012).
- [21] Magne Jorgensen and Martin Shepperd. 2007. A systematic review of software development cost estimation studies. *IEEE Transactions on Software Engineering* (2007).
- [22] Eirini Kalliamvakou, Georgios Gousios, Kelly Blincoe, Leif Singer, Daniel M. German, and Daniela Damian. 2014. The promises and perils of mining GitHub. In *Proc. MSR*.
- [23] Yasutaka Kamei, Emad Shihab, Bram Adams, Ahmed E. Hassan, Audris Mockus, Anand Sinha, and Naoyasu Ubayashi. 2013. A large-scale empirical study of just-in-time quality assurance. In *Proc. ICSE*.
- [24] Sunghun Kim, E. James Whitehead, and Yi Zhang. 2008. Classifying software changes: Clean or buggy?. In *Proc. TSE*.
- [25] Carlene Lebeuf, Margaret-Anne Storey, and Alexey Zagalsky. 2018. Software bots. *IEEE Software* (2018).

- 581 [26] Hao Li, Haixiang Zhang, and Ahmed E. Hassan. 2025. The Rise of AI Team- 639  
 582 mates in Software Engineering (SE) 3.0: How Autonomous Coding Agents Are 640  
 583 Reshaping Software Engineering. *arXiv preprint arXiv:2507.15003* (2025). 641  
 584 [27] Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, 642  
 Rémi Leblond, Tom Eccles, et al. 2022. Competition-Level Code Generation with 643  
 AlphaCode. *Science* (2022). 644  
 585 [28] Shane McIntosh, Yasutaka Kamei, Bram Adams, and Ahmed E. Hassan. 2014. 645  
 586 The impact of code review coverage and code review participation on software 646  
 587 quality evaluations. In *Proc. MSR*. 647  
 588 [29] Audris Mockus, Roy T. Fielding, and James D. Herbsleb. 2000. A case study of 648  
 589 open source software development: The Apache server. In *Proc. ICSE*. 649  
 590 [30] Peng et al. 2023. The Impact of AI on Developer Productivity: Evidence from 650  
 591 GitHub Copilot. *arXiv:2302.06590* (2023). 651  
 592 [31] Foyzur Rahman, Asif Partho, and David Meder. 2014. Characterizing the toxic 652  
 593 code in large software systems. In *Proc. ICSE*. 653  
 594 [32] Foyzur Rahman, Daryl Posnett, and Premkumar Devanbu. 2012. Recalling the 654  
 595 imprecision of cross-project defect prediction. In *Proc. FSE*. 655  
 596 [33] Baishakhi Ray, Daryl Posnett, Vladimir Filkov, and Premkumar Devanbu. 2014. 656  
 597 A large scale study of programming languages and code quality in github. In *Proc. FSE*. 657  
 598 [34] Rigby and Bird. 2013. Convergent Contemporary Software Peer Review Practices. 658  
 599 In *Proc. ESEC/FSE*. 659  
 600 [35] Igor Steinmacher, Marco Aurélio Graciotto Silva, Marco Aurélio Gerosa, and 660  
 601 David F. Redmiles. 2015. Social Barriers Faced by Newcomers Placing Their 661  
 602 First Contribution in Open Source Software Projects. In *Proc. CSCW*. 662  
 603 [36] Margaret-Anne Storey and Alexey Zagalsky. 2016. The who, what, how, and 663  
 604 why of software bots. In *Proc. ICSE*. 664  
 605 [37] Thongtanunam et al. 2017. Review Participation in Modern Code Review. *Em- 665  
 606 pirical Software Engineering* (2017). 666  
 607 608 609 610 611 612 613 614 615 616 617 618 619 620 621 622 623 624 625 626 627 628 629 630 631 632 633 634 635 636 637 638
- [38] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozire, et al. 2023. LLaMA: Open and Efficient Foundation Language Models. *arXiv preprint arXiv:2302.13971* (2023). 639  
[39] Christoph Treude, Ohad Barzilay, and Margaret-Anne Storey. 2011. How do programmers ask and answer questions on the web?. In *Proc. ICSE*. 640  
[40] Jason Tsay, Laura Dabbish, and James Herbsleb. 2014. Influence of Social and Technical Factors for Evaluating Contribution in GitHub. In *Proc. ICSE*. 356–366. 641  
[41] Priyan Vaithilingam, Tianyi Zhang, and Elena L. Glassman. 2022. Expectation vs. Experience: Evaluating the Usability of Code Generation Tools. In *Proc. CHI*. 642  
[42] Bogdan Vasilescu et al. 2015. Quality and Productivity Outcomes Relating to Continuous Integration in GitHub. *Proc. ESEC/FSE* (2015), 805–816. 643  
[43] Bogdan Vasilescu, Daryl Posnett, Baishakhi Ray, Mark G.J. van den Brand, Alexander Serebrenik, Premkumar Devanbu, and Vladimir Filkov. 2015. Gender and tenure diversity in GitHub teams. In *Proc. CHI*. 644  
[44] Wessel et al. 2018. The Power of Bots in OSS Projects. In *Proc. CSCW*. 645  
[45] Mairieli Wessel, Bruno M. de Souza, Igor Steinmacher, Igor S. Wiese, Ivanilton Polato, Ana Paula Chaves, and Marco A. Gerosa. 2019. How bots help developers on GitHub. In *Proc. ICSE*. 646  
[46] Mairieli Wessel, Alexander Serebrenik, Igor Wiese, Igor Steinmacher, and Marco A. Gerosa. 2021. The Effects of Adopting Code Review Bots on Pull Requests to OSS Projects. In *Proc. ICSME*. 647  
[47] Yue Yu et al. 2015. Wait For It: Determinants of Pull Request Latency in GitHub. 648  
 In *Proc. MSR*. 649  
[48] Yue Yu, Huaimin Wang, Gang Yin, and Charles X. Ling. 2016. Reviewer recommendation for pull-requests in GitHub: What can we learn from code review assignment?. In *Proc. ICSE*. 650  
[49] Albert Ziegler, Eirini Kalliamvakou, X. Alice Li, Andrew Rice, Devon Rifkin, 651  
 Shawn Simister, Ganesh Sitnikov, and Edward Tehrani. 2022. Productivity assessment of neural code generation. *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming* (2022). 652  
 653  
 654  
 655  
 656  
 657  
 658  
 659  
 660  
 661  
 662  
 663  
 664  
 665  
 666  
 667  
 668  
 669  
 670  
 671  
 672  
 673  
 674  
 675  
 676  
 677  
 678  
 679  
 680  
 681  
 682  
 683  
 684  
 685  
 686  
 687  
 688  
 689  
 690  
 691  
 692  
 693  
 694  
 695